

# CS3211 Cheatsheet AY22/23 Sem 2

by Richard Willie

## Concurrent Programming

### Concurrency and Parallelism

- **Concurrency:**

- Two or more tasks can start, run, and complete in overlapping time periods.
- They might not be running (executing on the CPU) at the same time.
- Two or more execution flows make *progress* at the same time by interleaving their executions or by executing instructions (on CPU) at exactly the same time

- **Parallelism:**

- Two or more tasks can run (execute) simultaneously, at the exact same time.
- Tasks do not only make progress, they actually execute at *simultaneously*.

### Race Condition

- **Criteria:**

1. Two concurrent threads access a shared resources without any synchronization.
2. At least one thread modifies the shared resources.

- A solution to the critical-section problem must satisfy the following three requirements:

1. **Mutual exclusion**

- If one thread is in the critical section, than no other is.

2. **Progress**

- If thread  $T$  is not in the critical section, then  $T$  cannot prevent other threads from entering the critical section.
- A thread in the critical section will eventually leave it.

3. **Bounded waiting (no starvation)**

- There exists a bound on the number of times that other threads are allowed to enter the critical section after a thread has made a request to enter the critical section.
- In other words, if thread  $T$  is waiting to enter its critical section, then  $T$  will eventually enter the critical section.

### Deadlock

- We say that a set of processes is in a deadlocked state when every processes in the set is waiting for an event that can be caused only by another process in the set.
- Deadlock can exist if and only if the following four conditions hold simultaneously:
  1. **Mutual exclusion** – at least one resource must be held in a non-sharable mode.
  2. **Hold and wait** – there must be one process holding one resource and waiting for another resource.
  3. **No preemption** – resources cannot be preempted, i.e. critical sections cannot be aborted externally.
  4. **Circular wait** – there must exist a set of processes  $\{P_1, P_2, \dots, P_n\}$  such that  $P_1$  is waiting for  $P_2$ ,  $P_2$  for  $P_3$ , etc.

### Starvation

- Starvation is a situation where a process is prevented from making progress because some other process has the resource it requires.
- Starvation is a side effect of the scheduling algorithm:
  - A high priority process always prevents a low priority process from running on the CPU.
  - One thread always beats another when acquiring a lock.

### Blocking Algorithm

- Algorithms that use mutexes, condition variables, and futures to synchronize data are called *blocking algorithms*.
  - Note: A spinlock is non-blocking, as it spins until the `test_and_set` is successful.

### Non-blocking Algorithm

- An algorithm is called *non-blocking* if failure or suspension of any thread cannot cause failure or suspension of another thread.
- Disadvantages of locks:
  - **Contention:** Some threads/processes have to wait until a lock (or a whole set of locks) is released. If one of the threads holding a lock dies, stalls, blocks, or enters an infinite loop, other threads waiting for the lock may wait indefinitely until the computer is power cycled.
  - **Priority inversion:** A low-priority thread/process holding a common lock can prevent high-priority threads/processes from proceeding.

- **Convoying:** All other threads have to wait if a thread holding a lock is descheduled due to a time-slice interrupt or page fault.
- **Composability:** It is hard to combine small, correct lock-based modules into equally correct larger programs without modifying the modules or at least knowing about their internals

- **Wait-freedom:**

- Wait-freedom is the strongest non-blocking guarantee of progress, combining guaranteed system-wide throughput with starvation-freedom.
- An algorithm is wait-free if every operation has a bound on the number of steps the algorithm will take before the operation completes.

- **Lock-freedom:**

- Lock-freedom allows individual threads to starve but guarantees system-wide throughput.
- An algorithm is lock-free if, when the program threads are run for a sufficiently long time, at least one of the threads makes progress (for some sensible definition of progress).
  - \* In particular, if one thread is suspended, then a lock-free algorithm guarantees that the remaining threads can still make progress.
  - \* Hence, if two threads can contend for the same mutex lock or spinlock, then the algorithm is *not* lock-free.
- All wait-free algorithms are lock-free.

- **Obstruction-freedom:**

- Obstruction-freedom is the weakest natural non-blocking progress guarantee.
- An algorithm is obstruction-free if at any point, a single thread executed in isolation (i.e., with all obstructing threads suspended) for a bounded number of steps will complete its operation.
- All lock-free algorithms are obstruction-free.

## C++

### Memory Model

#### Modification Orders

- Every object in C++ has a *modification order* composed of all the writes to that object from all threads in the program.
- Modification order varies between runs, but in any given execution of the program, we must ensure that all threads agree on the modification order.
- This requirements means that:

- Once a thread has seen a particular entry in the modification order,
  - \* subsequent reads from that thread must return later values, and
  - \* subsequent writes from that thread to that object must occur later in the modification order.
- A read of an object that follows a write to that object in the same thread must either return the value written or another value that occurs later in the modification order of that object.
- Although all threads must agree on the modification orders of each individual object in a program, they don't necessarily have to agree on the relative order of operations on separate objects.

### Sequentially-consistent Ordering

- If all operations on instances of atomic types are sequentially consistent, the behavior of the multithreaded program is *as if* all these operations were performed in some particular sequence by a single thread.
- In other words, all threads must see the same order of operations.
- We can reason about sequentially consistent ordering with the following mental model:
  - There's a single global order of events that is agreed by all threads.
  - The operations from different threads are neatly interleaved one after another.
  - Once a side-effect is visible to a thread, it is also visible to all other threads.

### Non-sequentially Consistent Ordering

- For non-sequentially consistent memory orderings, we *cannot* use the same mental model as the one we used for sequential consistency. This means that:
  - Operations from different threads don't neatly interleaved one after another.
  - Threads don't have to agree on the order of events (operations).
- In the absence of other ordering constraints, the only requirement is that all threads agree on the modification order of each individual variable.

## Acquire-release ordering

- Under this ordering model, atomic loads are *acquire* operations (`std::memory_order_acquire`), and atomic stores are *release* operations (`std::memory_order_release`).
- If an atomic store in thread *A* is tagged `memory_order_release`, an atomic load in thread *B* from the same variable is tagged `memory_order_acquire`, and the load in thread *B* reads a value written by the store in thread *A*, *then* the store in thread *A* *synchronizes-with* the load in thread *B*.
- All memory writes (non-atomic and relaxed atomic) that *happened-before* the atomic store from the point of view of thread *A*, become visible side-effects in thread *B*.
  - That is, once the atomic load is completed, thread *B* is guaranteed to see everything thread *A* wrote to memory.
  - This promise only holds if *B* actually returns the value that *A* stored, or a value from later in the release sequence.
- The synchronization is established only between the threads *releasing* and *acquiring* the same atomic variable. Other threads can see different order of memory accesses than either or both of the synchronized threads.

## Relaxed Ordering

- Operations on atomic types performed with relaxed ordering don't participate in *synchronizes-with* relationships.
- Operations on the same variable on within a single thread still obey *happens-before* relationship, but there's almost no requirement on ordering relative to other threads.

## Reordering Constraints on `std::memory_order`

- `std::memory_order_relaxed`: No reordering constraints.
- `std::memory_order_acquire`: No reads or writes in the current thread can be reordered *before* this load.
- `std::memory_order_release`: No reads or writes in the current thread can be reordered *after* this store.
- `std::memory_order_acq_rel`: No memory reads or writes in the current thread can be reordered *before* the load, nor *after* the store.

## Fences

- `std::atomic_thread_fence` establishes memory synchronization ordering for non-atomic and relaxed atomic accesses.
- Note however, that at least one atomic operation is required to set up the synchronization, as described below.
- **Fence-atomic synchronization:**
  - A release fence *F* in thread *A* *synchronizes-with* atomic *acquire operation* *Y* in thread *B*, if
    - \* there exists an atomic store *X* (with any memory order),
    - \* *Y* reads the value written by *X* (or the value would be written by release sequence headed by *X* if *X* was a release operation), and
    - \* *F* is *sequenced-before* *X* in thread *A*.
  - In this case, all non-atomic and relaxed atomic stores that are *sequenced-before* *F* in thread *A* will *happen-before* all non-atomic and relaxed atomic loads from the same locations made in thread *B* after *Y*.
- **Atomic-fence synchronization:**
  - An atomic *release operation* *X* in thread *A* *synchronizes-with* an acquire fence *F* in thread *B*, if
    - \* there exists an atomic read *Y* (with any memory order),
    - \* *Y* reads the value written by *X* (or by the release sequence headed by *X*), and
    - \* *Y* is *sequenced-before* *F* in thread *B*.
  - In this case, all non-atomic and relaxed non-atomic stores that are *sequenced-before* *X* in thread *A* will *happen-before* all non-atomic and relaxed atomic loads from the same locations made in thread *B* after *F*.
- **Fence-fence synchronization:**
  - A release fence *FA* in thread *A* *synchronizes-with* an acquire fence *FB* in thread *B*, if
    - \* there exists an atomic object *M*,
    - \* there exists an atomic write *X* (with any memory order) that modifies *M* in thread *A*,
    - \* *FA* is *sequenced-before* *X* in thread *A*,
    - \* there exists an atomic read *Y* (with any memory order) in thread *B*,
    - \* *Y* reads the value written by *X* (or the value would be written by release sequence headed by *X* if *X* were a release operation), and
    - \* *Y* is *sequenced-before* *FB* in thread *B*.

- In this case, all non-atomic and relaxed atomic stores that are sequenced-before *FA* in thread *A* will *happen-before* all non-atomic and relaxed atomic loads from the same locations made in thread *B* after *FB*.

## Debugging Tools for Multithreaded Programs

### Valgrind

- Valgrind is an instrumentation framework for building dynamic analysis tools.

- **Memcheck:**

- Memcheck detects memory-management problems, and is aimed primarily at C and C++ programs.
- Memcheck implementation uses *shadow memory*, i.e. all reads and writes of memory are checked, and calls to malloc/new/free/delete are intercepted.
- As a result, Memcheck can detect if your program:
  - \* Accesses memory it shouldn't.
  - \* Uses uninitialised values in dangerous ways.
  - \* Leaks memory.
  - \* Does bad frees of heap blocks (double frees, mismatched frees).
  - \* Passes overlapping source and destination memory blocks to memcpy() and related functions.
- Memcheck runs programs about 10-30x slower than normal.

- **Helgrind:**

- Helgrind is a Valgrind tool for detecting synchronisation errors in C and C++ programs that use the POSIX pthreads threading primitives.
- Helgrind can detect three classes of errors:
  - \* Misuses of the POSIX pthreads API.
  - \* Potential deadlocks arising from lock ordering problems.
  - \* Data races – accessing memory without adequate locking or synchronisation.
- Helgrind runs programs about 100x slower than normal.
- Deadlock detection:
  - \* Helgrind builds a directed graph indicating the order in which locks have been acquired.
- Data race detection:
  - \* Builds a directed acyclic graph representing the collective happens-before dependencies.
  - \* Monitors all memory accesses.

## Sanitizers

- **AddressSanitizer:**

- AddressSanitizer can detect the following types of bugs:
  - \* Out-of-bounds accesses to heap, stack and globals
  - \* Use-after-free
  - \* Use-after-return
  - \* Use-after-scope
  - \* Double-free, invalid free
  - \* Memory leaks
- AddressSanitizer uses a shadow memory scheme to detect memory bugs.
- AddressSanitizer does not detect any uninitialized memory reads (this is detected by MemorySanitizer).
- AddressSanitizer is also not capable of detecting all arbitrary memory corruption bugs, nor all arbitrary write bugs due to integer underflow/overflows.
- Typical slowdown introduced by AddressSanitizer is 2x.

- **ThreadSanitizer:**

- ThreadSanitizer is a tool that detects data races and deadlocks.
- Typical slowdown introduced by ThreadSanitizer is about 5-15x.
- Typical memory overhead introduced by ThreadSanitizer is about 5-10x.

## Go

### Goroutines

- A *goroutine* is a lightweight thread managed by the Go runtime.
- Goroutines run in the same address space, so access to shared memory must be synchronized.
- Goroutines are not garbage collected.

## Channels

- Philosophy of Go: “Do not communicate by sharing memory. Instead, share memory by communicating.”
- Channels are a typed conduit through which you can send and receive values.

```
1  ch <- v    // Send v to channel ch
2  v := <-ch  // Receive from ch, and assign value to v
```

### • Unbuffered channel:

```
1  ch := make(chan int)
```

- Sends and receives block until the other side is ready.

### • Buffered channel:

```
1  ch := make(chan int, 100)
```

- Sends to a buffered channel block only when the buffer is full.
- Receives block when the buffer is empty.

### • Range and Close:

- A sender can **close** a channel to indicate that no more values will be sent.
- Receivers can test whether a channel has been closed by assigning a second parameter to the receive expression:

```
1  v, ok := <-ch
```

**ok** is **false** if there are no more values to receive and the channel is closed.

- The loop receives values from the channel repeatedly until it is closed.

```
1  ch := make(chan int, 10)
2  ...
3  for i := range c {
4      ...
5  }
```

### • Ownership of a channel:

- Owner should:
  - \* Instantiate the channel.
  - \* Have write-access view into the channel. (**chan** or **chan<-**)

- \* Close the channel.
- \* Pass ownership to another goroutine.
- \* Encapsulate the channel and expose to consumers via a read-only view into the channel. (**<-chan**)

### – Consumer should:

- \* Know when a channel is closed.
- \* Responsibly handle blocking for any reason.

### • Select:

- The **select** statement lets a goroutine wait on multiple communication operations.
- A select **blocks** until one of its cases can run, then it executes that case. It chooses one at random if multiple are ready.

```
1  for {
2      select {
3          case i := <-c:
4              // use i
5          case <-quit:
6              ...
7          return
8          default:
9              ...
10     }
11 }
```

## Rust

### Ownership, Borrowing, and Lifetimes

#### • Ownership rules:

- Each value in Rust has an **owner**.
- There can only be one owner at a time.
- When the owner goes out of scope, the value will be dropped.

#### • The ownership rules prevent **double free**.

#### • The rules of references:

- At any given time, you can have **either** one mutable reference or any number of immutable references. This rule prevents data race.
- Every reference in Rust has a **lifetime**, which is the scope for which that reference is valid. This rule prevents **dangling references**.

- The Rust compiler has a **borrow checker** that compares scopes to determine whether all borrows are valid.

## Smart Pointers in Rust

- Smart pointers are usually implemented using structs.
- Unlike an ordinary struct, smart pointers implement the `Deref` and `Drop` traits.
  - The `Deref` trait allows you to customize the behavior of the *dereference operator*.
  - The `Drop` trait allows you to customize the code that's run when an instance of the smart pointer goes out of scope. (e.g. RAII)

## Box

- Boxes don't have performance overhead.
- Use cases:
  - When you have a type whose size can't be known at compile time, and you want to use a value of that type in a context that requires an exact size.
  - When you have a large amount of data and you want to transfer ownership but ensure the data won't be copied when you do so.
  - When you want to own a value and you care only that it's a type that implements a particular trait rather than being of a specific type.
- Enabling recursive types with Boxes:
  - Recursive types pose an issue because at compile time Rust needs to know how much space a type takes up.
  - Implementing a *cons list*:

```
1 enum List {
2     Cons(i32, Box<List>),
3     Nil,
4 }
```

## Rc, the Referenced Counted Smart Pointer

- To enable multiple ownership explicitly, we use `Rc<T>`, which is an abbreviation for *reference counting*.
  - The `Rc<T>` type keeps track of the number of references to a value to determine whether or not the value is still in use.
  - If there are zero references to a value, the value can be cleaned up without any references becoming invalid.

```
1 enum List {
2     Cons(i32, Rc<List>),
3     Nil,
4 }
5
6 fn main() {
7     let a = Rc::new(Cons(5, Rc::new(Cons(10, Rc::new(Nil)))));
8     let b = Cons(3, Rc::clone(&a));
9     let c = Cons(4, Rc::clone(&a));
10 }
```

- Cloning an `Rc<T>` increases the reference count.
  - We can get the reference count by calling the `Rc::strong_count` function.
  - The function is named `strong_count` rather than `count` because `Rc<T>` type also has a `weak_count`.
- `Rc<T>` allows us to share data between multiple parts of our program for reading only.

## RefCell and the Interior Mutability Pattern

- *Interior mutability* is a design pattern in Rust that allows you to mutate data even when there are immutable references to that data.
  - To mutate data, the pattern uses `unsafe` code.
  - Unsafe code indicates to the compiler that we're checking the rules manually instead of relying on the compiler to check them for us.
- Similar to `Rc<T>`, `RefCell<T>` is only for use in single-threaded scenarios.
- Here is a recap of the reasons to choose `Box<T>`, `Rc<T>`, or `RefCell<T>`:
  - `Rc<T>` enables multiple owners of the same data; `Box<T>` and `RefCell<T>` have single owners.
  - `Box<T>` allows immutable or mutable borrows checked at compile time; `Rc<T>` allows only immutable borrows checked at compile time; `RefCell<T>` allows immutable or mutable borrows checked at runtime.
  - Because `RefCell<T>` allows mutable borrows checked at runtime, you can mutate the value inside the `RefCell<T>` even when the `RefCell<T>` is immutable.
- With `RefCell<T>`, we use the `borrow` and `borrow_mut` methods.
  - The `borrow` method returns the smart pointer type `Ref<T>`.
  - The `borrow_mut` method returns the smart pointer type `RefMut<T>`.

- The `RefCell<T>` keeps track of how many `Ref<T>` and `RefMut<T>` smart pointers are currently active.
  - Just like the compile-time borrowing rules, `RefCell<T>` lets us have many immutable borrows or one mutable borrow at any point in time.
- We can have multiple owners of mutable data by combining `Rc<T>` and `RefCell<T>`.

### Reference Cycles Can Leak Memory

- Rust's memory safety guarantee makes it difficult, but not impossible, to accidentally create memory that is never cleaned up (known as *memory leak*).
- Preventing memory leaks entirely is not one of Rust's guarantees, meaning memory leaks are memory safe in Rust.
- We can see that Rust allows memory leaks by using `Rc<T>` and `RefCell<T>`: it's possible to create references where items refer to each other in a cycle.

### Fearless Concurrency

- By leveraging ownership and type checking, many concurrency errors are compile-time errors in Rust rather than runtime errors.
- The Rust standard library uses a 1:1 model of thread implementation, whereby a program uses one operating system thread per one language thread.

### Using Threads in Rust

```
1 fn main() {
2     let handle = thread::spawn(|| { ... });
3     handle.join().unwrap();
4 }
```

- The return type of `thread::spawn` is `JoinHandle`.
- A `JoinHandle` is an owned value that, when we call the `join` method on it, will wait for its thread to finish.
- We use the `move` keyword to force closure to take ownership of the values it uses from the environment.

```
1 fn main() {
2     let v = vec![1, 2, 3];
3     let handle = thread::spawn(move || {
4         println!("{:?}", v);
5     });
6     handle.join().unwrap();
7 }
```

### Message-Passing Concurrency in Rust

- Rust also supports message-passing concurrency with the use of *channels*.

```
1 fn main() {
2     let (tx, rx) = mpsc::channel();
3     thread::spawn(move || {
4         let val = String::from("hi");
5         tx.send(val).unwrap();
6     });
7     let received = rx.recv().unwrap();
8     println!("Got: {}", received);
9 }
```

- A channel has two halves: a transmitter and a receiver.
- A channel is said to be *closed* if either the transmitter or receiver half is dropped.

- The `send` function:

- takes ownership of its parameter, and when the value is moved, the receiver takes ownership of it.
- returns a `Result<T, E>` type, so if the receiver has already been dropped and there's nowhere to send a value, the send operation will return an error.

- The receiver has two useful methods: `recv` and `try_recv`.

- `recv` will block the thread's execution and wait until a value is sent down the channel.
- `try_recv` method doesn't block, but will instead return a `Result<T, E>` immediately: an `Ok` value holding a message if one is available and an `Err` value if there aren't any messages this time.

- We can treat `rx` as an iterator when we want to receive multiple messages.

```
1 for received in rx { ... }
```

- When all transmitters have been dropped, the channel is closed, and the iteration will end.

- `mpsc` is an acronym for *multiple producer, single consumer*.

- We can create multiple transmitters by cloning the transmitter.



## Shared-State Concurrency in Rust

- Mutexes in Rust implement a strategy called *poisoning* where a mutex is considered poisoned whenever a thread panics while holding the mutex.
  - For a mutex, this means that the `lock` and `try_lock` methods return a `Result` which indicates whether a mutex has been poisoned or not.
  - A poisoned mutex, however, does not prevent all access to the underlying data. The `PoisonError` type has an `into_inner` method which return the guard to allow access to data.
- `Mutex<T>` is a smart pointer.
  - The call to `lock` returns a smart pointer called `MutexGuard`, wrapped in a `LockResult` that we handled with the call to `unwrap`.
  - The `MutexGuard` smart pointer implements `Deref` to point to our inner data.
  - The smart pointer also has a `Drop` implementation that releases the lock automatically when a `MutexGuard` goes out of scope. (RAII)
- To share a `Mutex<T>` between multiple threads, we wrap it with `Arc<T>`, which is an *atomically reference counted smart pointer*.

```
1 fn main() {
2     let counter = Arc::new(Mutex::new(0));
3     let mut handles = vec![];
4     for _ in 0..10 {
5         let counter = Arc::clone(&counter);
6         let handle = thread::spawn(move || {
7             let mut num = counter.lock().unwrap();
8             *num += 1;
9         });
10    handles.push(handle);
11 }
12 for handle in handles {
13     handle.join().unwrap();
14 }
15 println!("Result: {}", *counter.lock().unwrap());
16 }
```

- Notice that `counter` is immutable but we could get a mutable reference to the value inside it; this means that `Mutex<T>` provides interior mutability, as the `Cell` family does.
- Using `Mutex<T>` does not guarantee that we are safe from deadlocks.

## Sync and Send Traits

- Allowing transference of ownership between threads:
  - The `Send` marker trait indicates that the ownership of values of the type implementing `Send` can be transferred between threads.
  - Almost every Rust type in `Send`, but there are some exceptions, including `Rc<T>`.
  - Any type composed entirely of `Send` types is automatically marked as `Send` as well.
  - Almost all primitive types are `Send`, aside from raw pointers.
- Allowing access from multiple threads:
  - The `Sync` marker trait indicates that it is safe for the type implementing `Send` to be referenced from multiple threads.
  - Any type `T` is `Sync` if `&T` is `Send`, meaning the reference can be sent safely to another threads.
  - Primitive types are `Sync`, and types composed entirely of types that are `Sync` are also `Sync`.
  - The `RefCell<T>` type and the family of related `Cell<T>` types are not `Sync`. The implementation of borrow-checking that `RefCell<T>` does at runtime is not thread-safe.
  - The smart pointer `Mutex<T>` is `Sync`.
- Manually implementing the `Send` and `Sync` traits involves implementing `unsafe` Rust code.

## Asynchronous Programming in Rust

### Futures

- The problem with threads:
  - Context switching cost: Each switch is expensive.
  - Memory overhead: Each thread has its own stack space that needs to get managed by the OS.
- Non-blocking I/O:
  - We could have `read()` return a special error value instead of blocking.
  - If we see that a client hasn't sent us anything yet, we can do other useful work on this thread e.g. reading from other descriptors we're managing.
  - `epoll` is a kernel-provided mechanism that notifies us of what fds are ready for I/O.
  - This allows us to have concurrent I/O with one thread!
- State management:
  - Non-blocking I/O code looks okay in theory, but in reality we need to figure out how to manage the state associated with each conversation.



- *Futures* in Rust allows us to keep track of in-progress operations along with associated state, in one package.
- Rust docs: Futures are single eventual values produced by asynchronous computations.

- The **Future** trait:

```

1  trait Future {
2      type Output;
3      fn poll(&mut self, cx: &mut Context) -> Poll<Self::Output>;
4  }
5
6  enum Poll<T> {
7      Ready(T),
8      Pending,
9  }

```

- The executor thread should call `poll()` on the future to start it off.
- It will run code until it can no longer progress.
  - \* If the future is complete, returns `Poll::Ready(T)`.
  - \* If future needs to wait for some event, returns `Poll::Pending`, and allows the single thread to work on another futures.
  - \* When `poll()` is called, `Context` structure passed in.
  - \* Includes a `wake()` function that is set to be called when future can make progress again. (This is implemented internally using system calls)
  - \* After `wake()` called, executor can use `Context` to see which Future can be polled to make new progress.

- How executors work:

- An executor loops over futures that can currently make progress, calling `poll()` on them to give them attention until they need to wait again.
  - \* When no futures can make progress, the executor goes to sleep until one or more futures calls `wake()`.
  - \* Once awakened, the executor goes through those futures, `poll()`ing them.

- Futures should not block:

- If code within a future causes the thread to sleep, the executor running that code is going to sleep!
- This defeats the purpose of the system since the executor cannot continue to other futures.
- Asynchronous code needs to use non-blocking versions of *everything*.

## Async/await

- Working with futures is not ergonomic:

- Futures are composed with various combinators.
- Chaining combinators may result in strange decomposition, which is bad for abstraction.

- Rust introduced **async/await** which are syntactic sugar to eliminate the issues.

- An **async** function is a function that returns a **Future**.

- **.await** waits for a future and gets its value.

- You can only use **await** in **async** functions.

```

1  async fn add_to_inbox(email_id: u64, recipient_id: u64)
2      -> Result<(), Error> {
3      let message = load_message(email_id).await?;
4      let recipient = get_recipient(recipient_id).await?;
5      recipient.verify_has_space(&message)?;
6      recipient.add_to_inbox(message).await
7  }

```

- The code gets compiled into a **Future** with `poll()` method.
- There are 5 places where the executor might be paused, or not actively executing. Thus internally, we can use an **enum** to store the state for these possibilities.

- Implications:

- **async** functions have no stack! (sometimes called *stackless coroutines*)
  - \* The executor thread still has a stack, but it isn't used to store state when switching between **async** tasks. All state is self contained in the generated **Future**.
- No recursion! The **Future** returned by an **async** function needs to have a fixed size known at compile time.
- **async** functions are one of Rust's *zero-cost abstractions*, meaning there is no runtime overhead.

- Async code makes sense when...

- You need an extremely high degree of concurrency.
- Work is primarily I/O bound.
  - \* Context switching overhead is expensive only if you're using a tiny fraction of the time slice.
  - \* If you're doing a lot of work on the CPU for an extended period of time, you might prevent the executor from running other tasks.

## Concurrency Libraries in Rust

### Tokio

- At a high level, Tokio provides a few major components:
  - A multi-threaded runtime for executing asynchronous code.
  - An asynchronous version of the standard library.
- When not to use Tokio:
  - Speeding up CPU-bound computations by running them in parallel on several threads. Tokio is designed for I/O-bound applications. You should be using *rayon* instead.
  - Reading a lot of files, as Tokio provides no advantage here compared to an ordinary threadpool.
  - Sending a single web request. The place where Tokio gives you an advantage is when you need to do many things at the same time.

### Rayon

- A data-parallelism library for Rust.
- It is extremely lightweight and makes it easy to convert a sequential computation into a parallel one.
- It also guarantees data-race freedom.

### Crossbeam

- Provides an implementation of scoped threads.
  - Scoped threads are allowed to borrow variables on stack, since its mechanism guarantee to the compiler that spawned threads will be joined before the scope ends.
  - In other words, if a variable is borrowed by a thread inside the scope, the thread must complete before the variable is destroyed.
- Provides multi-producer multi-consumer channels for message passing.
- Provides exponential backoff utility.
  - Performs exponential backoff in spin loops.
  - Rationale: It's not helpful to simply retry as fast as possible. If the resource is overloaded, requesting it more will make it even more overloaded!

## Classical Synchronization Problems

### Readers-writers

#### Solution 1

- C++ implementation:

```
1  std::mutex mut;
2
3  void reader() {
4      std::shared_lock lock{mut};
5      // Critical section
6  }
7
8  void writer() {
9      std::unique_lock lock{mut};
10     // Critical section
11 }
```

- Issue: Starvation of writers is possible.

#### Solution 2 (Starvation-free)

- C++ implementation:

```
1  std::mutex mut_write;
2  std::mutex mut_read;
3
4  void reader() {
5      { std::scoped_lock lock1{mut_write}; }
6      std::shared_lock lock2{mut_read};
7      // Critical section
8  }
9
10 void writer() {
11     std::scoped_lock lock1{mut_write};
12     std::unique_lock lock2{mut_read};
13     // Critical section
14 }
```

- Issue: Starvation of writers is possible.

### Barrier

#### Solution

- C++ implementation:

```

1  class Barrier {
2  private:
3      std::ptrdiff_t expected;
4      std::ptrdiff_t count;
5      std::mutex mut;
6      std::counting_semaphore<> turnstile1;
7      std::counting_semaphore<> turnstile2;
8
9  public:
10     Barrier(std::ptrdiff_t expected)
11         : expected{expected}, count{0}, mut{},
12           turnstile1{0}, turnstile2{1} {}
13
14     void arrive_and_wait() {
15         {
16             std::scoped_lock lock{mut};
17             count++;
18             if (count == expected) {
19                 turnstile2.acquire()
20                 turnstile1.release()
21             }
22         }
23
24         turnstile1.acquire();
25         turnstile1.release();
26
27         {
28             std::scoped_lock{mut};
29             count--;
30             if (count == 0) {
31                 // Reset the barrier
32                 turnstile1.acquire();
33                 turnstile2.release();
34             }
35         }
36
37         turnstile2.acquire();
38         turnstile2.release();
39     }
40 };

```

- Go implementation:

```

1  type Barrier struct {

```

```

2      wg1 sync.WaitGroup;
3      wg2 sync.WaitGroup;
4  }
5
6  func (b *Barrier) Init(expected int) {
7      b.wg1.Add(expected);
8      b.wg2.Add(expected);
9  }
10
11 func (b *Barrier) Wait() {
12     b.wg1.Done();
13     b.wg1.Wait();
14
15     // Reset the barrier
16     b.wg1.Add(1);
17
18     b.wg2.Done();
19     b.wg2.Wait();
20     b.wg2.Add(1);
21 }

```

## Dining Philosophers

```

1  def philosopher(i):
2      while True:
3          think()
4          get_chopsticks(i)
5          eat()
6          put_chopsticks(i)

```

## Solution 1

```

1  def get_chopsticks(i):
2      left_chopstick(i).lock()
3      right_chopstick(i).lock()
4
5  def put_chopsticks(i):
6      left_chopstick(i).unlock()
7      right_chopstick(i).unlock()

```

- Issue: Deadlock
- This solution is actually fine if we implement it in C++, since `std::scoped_lock` uses a deadlock avoidance algorithm when locking multiple mutexes.

## Solution 2

- Intuition: If we have one less philosophers at the table, deadlock can be avoided.
- We can control the number of philosophers at the table with a ‘footman’ semaphore.

```
1  footman = Semaphore(num_of_philosophers - 1)
2
3  def get_chopsticks(i):
4      footman.acquire();
5      left_chopstick(i).lock()
6      right_chopstick(i).lock()
7
8  def put_chopsticks(i):
9      footman.release();
10     left_chopstick(i).unlock()
11     right_chopstick(i).unlock()
```

- This solution is starvation-free if the `footman` semaphore has the following property:
  - If a thread is waiting at a semaphore, then the number of threads that will be woken before it is bounded.
- In C++, starvation is possible because mutexes and semaphores are not fair.

## Solution 3

- Intuition: To avoid deadlock, change the order in which the philosophers pick up the chopsticks.
- In Go implementation, we can use odd-even ring communication to avoid deadlock.

```
1  type Chopstick struct{}
2
3  type DiningTable struct {
4      numOfPhilosopher int
5      chopsticks []chan Chopstick
6  }
7
8  func (t *DiningTable) Init(numOfPhilosopher int) {
9      t.numOfPhilosopher = numOfPhilosopher
10     t.chopsticks = make([]chan Chopstick, 0, numOfPhilosopher)
11     for i := 0; i < numOfPhilosopher; i++ {
12         chopstick := make(chan chopstick, 1)
13         chopstick <- Chopstick{}
```

```
14         t.chopsticks = append(t.chopsticks, chopstick)
15     }
16 }
17
18 func (t *DiningTable) LeftChopstick(pid int) chan Chopstick {
19     return t.chopsticks[pid]
20 }
21
22 func (t *DiningTable) RightChopstick(pid int) chan Chopstick {
23     return t.chopsticks[(pid + 1) % t.numOfPhilosopher]
24 }
25
26 func (t *DiningTable) EvenChopstick(pid int) chan Chopstick {
27     if pid % 2 == 0 {
28         return t.LeftChopstick(pid)
29     } else {
30         return t.RightChopstick(pid)
31     }
32 }
33
34 func (t *DiningTable) OddChopstick(pid int) chan Chopstick {
35     if pid % 2 == 1 {
36         return t.LeftChopstick(pid)
37     } else {
38         return t.RightChopstick(pid)
39     }
40 }
41
42 func (t *DiningTable) Eat(pid int, eat_callback func(pid int)) {
43     evenChopstick := t.EvenChopstick(pid)
44     oddChopstick := t.OddChopstick(pid)
45
46     <-evenChopstick
47     <-oddChopstick
48
49     eat_callback(pid)
50
51     evenChopstick <- Chopstick{}
52     oddChopstick <- Chopstick{}
53 }
```

## Solution 4 (Tanenbaum's)

- For each philosopher there is a state variable that indicates whether the philosopher is thinking, eating, or waiting to eat ("hungry") and a semaphore that indicates whether the philosopher can start eating.

```

1  state = ["thinking" for i in range(n)]
2  sem = [Semaphore(0) for i in range(n)]
3  mutex = Semaphore(1)
4
5  def get_fork(i):
6      mutex.wait()
7      state[i] = "hungry"
8      test(i)
9      mutex.signal()
10     sem[i].wait()
11
12  def put_fork(i):
13      mutex.wait()
14      state[i] = "thinking"
15      test(left(i))
16      test(right(i))
17      mutex.signal()
18
19  def test(i):
20      if state[i] == "hungry" and
21         state[left(i)] == "eating" and
22         state[right(i)] == "eating":
23         state[i] == "eating"
24         sem[i].signal()

```

- Issue: Deadlock is not possible, but the solution is not starvation-free.

## Barbershop

### Solution

- C++ implementation:

```

1  template <size_t max_customer>
2  class Barbershop {
3  private:
4      size_t customer;
5      std::mutex mut;
6      std::counting_semaphore<> sem_customer;
7      std::counting_semaphore<> sem_barber;
8      std::counting_semaphore<> sem_customer_done;
9      std::counting_semaphore<> sem_barber_done;
10
11  public:
12      void barber() {
13          sem_customer.acquire();

```

```

14         sem_barber.release();
15         cut_hair();
16         sem_customer_done.acquire();
17         sem_barber_done.release();
18     }
19
20     void customer() {
21         {
22             std::scoped_lock lock{mut};
23             if (customer == max_customer) {
24                 balk();
25                 return;
26             }
27             customer++;
28         }
29
30         sem_customer.release();
31         sem_barber.wait();
32         get_hair_cut();
33         sem_customer_done.release();
34         sem_barber_done.acquire();
35
36         {
37             std::scoped_lock{mut};
38             customer--;
39         }
40     }
41 };

```

- Go implementation:

```

1  type Barbershop struct {
2      chairs chan int
3      barber chan struct{}
4  }
5
6  func (bs *Barbershop) Init(numOfChair int) {
7      bs.chairs = make(chan int, numOfChair)
8      bs.barber = make(chan struct{})
9  }
10
11  func (bs *Barbershop) Barber() {
12      for {
13          bs.barber <- struct{}{}

```

```

14     <-bs.chairs
15     cutHair()
16 }
17 }
18
19 func (bs *Barbershop) Customer() {
20     for {
21         select {
22             case bs.chairs <- 1:
23                 <-bs.barber
24                 getHairCut()
25             default:
26                 balk()
27         }
28     }
29 }

```

## H2O

### Solution

- C++ implementation:

```

1  class WaterFactory {
2  private:
3      std::barrier<> barrier;
4      std::counting_semaphore<> hydrogen_sem;
5      std::counting_semaphore<> oxygen_sem;
6
7  public:
8      WaterFactory()
9          : barrier{3}, hydrogen_sem{2}, oxygen_sem{1} {}
10
11     void hydrogen(void (*bond)()) {
12         hydrogen_sem.acquire();
13         barrier.arrive_and_wait();
14         bond();
15         hydrogen_sem.release();
16     }
17
18     void oxygen(void (*bond)()) {
19         oxygen_sem.acquire();
20         barrier.arrive_and_wait();
21         bond();
22         oxygen_sem.release();
23     }

```

```

24 };

```

- Go implementation using a daemon goroutine:

```

1  type WaterFactoryWithDaemon struct {
2      precomH chan chan struct{}
3      precomO chan chan struct{}
4  }
5
6  func NewFactoryWithDaemon() WaterFactoryWithDaemon {
7      wfd := WaterFactoryWithDaemon{
8          precomH: make(chan chan struct{}),
9          precomO: make(chan chan struct{}),
10     }
11
12     // Daemon goroutine
13     go func() {
14         h1 := <-wfd.precomH
15         h2 := <-wfd.precomH
16         o1 := <-wfd.precomH
17
18         h1 <- struct{}{}
19         h2 <- struct{}{}
20         o1 <- struct{}{}
21
22         <-h1
23         <-h2
24         <-o1
25     }()
26
27     return wfd
28 }
29
30 func (wfd *WaterFactoryWithDaemon) hydrogen(bond func()) {
31     commit := make(chan struct{})
32     wfd.precomH <- commit
33     <-commit
34     bond()
35     commit <- struct{}{}
36 }
37
38 func (wfd *WaterFactoryWithDaemon) oxygen(bond func()) {
39     commit := make(chan struct{})
40     wfd.precomO <- commit
41     <-commit

```

```

42     bond()
43     commit <- struct{}{}
44 }

```

- Downsides of the daemon-based approach:
  - The daemon is a bottleneck and a single point of failure.
  - The daemon goroutine does not go away and leaks memory.
- Go implementation using oxygen atoms as leader goroutines:

```

1  type WaterFactoryWithLeader struct {
2      oxygenMutex chan struct{}
3      precomH chan chan struct{}
4  }
5
6  func NewFactoryWithLeader() WaterFactoryWithLeader {
7      wf := WaterFactoryWithLeader{
8          oxygenMutex: make(chan struct{}), 1),
9          precomH: make(chan chan struct{}),
10     }
11
12     wf.oxygenMutex <- struct{}{}
13     return wf
14 }
15
16 func (wf *WaterFactoryWithLeader) hydrogen(bond func()) {
17     commit := make(chan struct{})
18     wf.precomH <- commit
19     <-commit
20     bond()
21     commit <- struct{}{}
22 }
23
24 func (wf *WaterFactoryWithLeader) oxygen(bond func()) {
25     <-wf.oxygenMutex
26
27     h1 := <-wf.precomH
28     h2 := <-wf.precomH
29
30     h1 <- struct{}{}
31     h2 <- struct{}{}
32
33     bond()
34
35     <-h1

```

```

36     <-h2
37
38     wf.oxygenMutex <- struct{}{}
39 }

```

- Rust implementation:

```

1  use tokio::sync::{mpsc, Barrier, Mutex, Semaphore};
2
3  async fn hydrogen(
4      id: usize, barrier: Arc<Barrier>,
5      sem: Arc<Semaphore>, chan: mpsc::Sender<usize>,
6  ) {
7      let _permit = sem.acquire().await.unwrap();
8      barrier.wait().await;
9      chan.send(id).await.unwrap();
10 }
11
12 async fn oxygen(
13     id: usize, barrier: Arc<Barrier>,
14     chan: Arc<Mutex<mpsc::Receiver<usize>>>
15 ) {
16     let mut chan_guard = chan.lock().await;
17     barrier.wait().await;
18     let h1 = chan_guard.recv().await.unwrap();
19     let h2 = chan_guard.recv().await.unwrap();
20     println!("H {} - O {} - H {}", h1, id, h2);
21 }
22
23 #[tokio::main]
24 async fn main() {
25     let barrier = Arc::new(Barrier::new(3));
26     let h_sem = Arc::new(Semaphore::new(2));
27     let (s, r) = mpsc::channel(2);
28     let r = Arc::new(Mutex::new(r));
29     let hydrogens = (0..200).map(|i| tokio::spawn(hydrogen(
30         i, barrier.clone(), h_sem.clone(), s.clone())));
31     let oxygens = (0..100).map(|i| tokio::spawn(oxygen(
32         i, barrier.clone(), r.clone())));
33     let join_handles = Iterator::chain(hydrogens, oxygens).collect::<
34         std::mem::drop(s);
35         std::mem::drop(r);
36         join_all(join_handles).await;
37 }

```